

Capstone_project_Olist

Carmen Marquez

8/12/2019

The goal of my capstone project is understand and find out a way to predict Olist's customers satisfaction. Olist(<https://olist.com/>) is a brazilian company that basically is a great sales channel it is present on the main marketplaces of Brazil and it is formed by thousands of retailers. It's something similar to the "brazilian Amazon". I got the data from kaggle (<https://www.kaggle.com/olistbr/brazilian-ecommerce> and <https://www.kaggle.com/olistbr/marketing-funnel-olist/home>). After having a look to all the files available I decided to include in my analysis only the most relevant ones for the customer satisfaction, since the memory of my pc could not handle the total number of observation of all the files.

So, the steps I followed to achieve my goal are the following: -First, I did a Exploratory Data Analysis to have a better understanding of the data I was dealing with -Second, I did a little of feature engineering, creating new variables and I clean the data dropping the variables that were not relevant or were missing to many values. Also I perform a sentiment analysis. -Third, I did some unsupervised analysis (Cluster analysis in particular) too see how homogeneous was the information and how it could be split into different categories. -Fourth, I did some supervised analysis. -1st. A multivariable lineal regression model using the review score as the dependent variable as a numeric variable. - 2n. An ordinal logistic regression model using the review score as the dependent variable as a factor -3rd. I run different machine learning models and chose the random forest as the best.

-Fifth, I decided to choose the ordinal logistic regression model and I interpret the results.

1.Imported and merged the data files

```
library(tidyverse)
```

```
## — Attaching packages —  
———— tidyverse 1.2.1 —
```

```
## ✓ ggplot2 3.2.1      ✓ purrr   0.3.2  
## ✓ tibble  2.1.3      ✓ dplyr   0.8.3  
## ✓ tidyr   0.8.3      ✓ stringr 1.4.0  
## ✓ readr   1.3.1      ✓ forcats 0.4.0
```

```
## — Conflicts —  
———— tidyverse_conflicts() —  
## * dplyr::filter() masks stats::filter()  
## * dplyr::lag()     masks stats::lag()
```

```
library(dplyr)
```

```
#We need to import the data
```

```
#As there are some blank values in our data, will replace it by NA
```

```
setwd("/home/achaparro/personal/Carmen")  
MQL <- read.csv("olist_marketing_qualified_leads_dataset.csv", na.strings = c("", "NA"))  
closed_deals <- read.csv("olist_closed_deals_dataset.csv", na.strings = c("", "NA"))  
total <- merge(closed_deals,MQL, by ="mql_id", all= TRUE)
```

```
order_items <- read.csv("olist_order_items_dataset.csv", na.strings = c("", "NA"))  
order_reviews <- read.csv("olist_order_reviews_dataset_EN.csv", na.strings = c("", "NA"))  
orders <- read.csv("olist_orders_dataset.csv", na.strings = c("", "NA"))
```

```
#The dataset "total" contains the Marketing funnel key variables, now let's merge that information with the  
Brazilian e-commerce public dataset
```

```
total1 <- merge(total,order_items, by ="seller_id", all= TRUE)
```

```
total2 <- merge(total1,orders, by ="order_id", all= TRUE)
```

```
total3 <- merge(total2,order_reviews, by ="order_id", all= TRUE)
```

2.Some Exploratory Data Analysis to understand better and clean our data

#1. We drop the identifier variables as they are not usefull for our purpose

```
remove02 = c("order_id", "seller_id", "mql_id", "sdr_id", "sr_id", "landing_page_id", "product_id", "customer_id", "review_id")
```

```
total4 = total3 %>% dplyr::select(-remove02)
```

#2. We calculate son new variables from the date variables

```
total4$won_date_new <- as.character(total4$won_date, format = "%Y-%m-%d")
total4$won_date_new <- as.Date(total4$won_date_new, format = "%Y-%m-%d")
total4$first_contact_date <- as.Date(total4$first_contact_date, format = "%Y-%m-%d")
total4$conversion_time <- (total4$won_date_new - total4$first_contact_date)
total4$conversion_time <- as.numeric(total4$conversion_time)

total4$order_delivered_customer_date <- as.Date(total4$order_delivered_customer_date, format = "%Y-%m-%d")
total4$order_purchase_timestamp <- as.Date(total4$order_purchase_timestamp, format = "%Y-%m-%d")
total4$delivery_time <- total4$order_delivered_customer_date - total4$order_purchase_timestamp
total4$delivery_time <- as.character(total4$delivery_time)
total4$delivery_time <- as.numeric(total4$delivery_time)
class(total4$delivery_time)
```

```
## [1] "numeric"
```

```
total4$review_creation_date <- as.Date(total4$review_creation_date, format = "%Y-%m-%d")
total4$feedback_time <- total4$review_creation_date - total4$order_purchase_timestamp
total4$feedback_time <- as.numeric(total4$feedback_time)

total4$order_estimated_delivery_date <- as.Date(total4$order_estimated_delivery_date, format = "%Y-%m-%d")
total4$delay_time <- total4$order_delivered_customer_date -
total4$order_estimated_delivery_date
total4$delay_time <- as.numeric(total4$delay_time)

total4$order_approved_at <- as.Date(total4$order_approved_at, format = "%Y-%m-%d")
total4$approval_time <- total4$order_approved_at - total4$order_purchase_timestamp
total4$approval_time <- as.numeric(total4$approval_time)
```

#3. We delete the date variables since we got the information we need from them in our new variables

```
remove03 = c("won_date", "shipping_limit_date", "order_purchase_timestamp", "order_approved_at", "order_delivered_carrier_date", "order_delivered_customer_date", "order_estimated_delivery_date", "review_creation_date", "review_answer_timestamp", "won_date_new", "first_contact_date")
```

```
total4 = total4 %>% dplyr::select(-remove03)
```

```
str(total4)
```

```
## 'data.frame': 121720 obs. of 23 variables:
## $ business_segment : Factor w/ 33 levels "air_conditioning",...: NA NA NA NA NA NA NA NA NA NA 6
...
## $ lead_type : Factor w/ 8 levels "industry","offline",...: NA NA NA NA NA NA NA NA NA NA
4 ...
## $ lead_behaviour_profile : Factor w/ 9 levels "cat","cat, wolf",...: NA NA NA NA NA NA NA NA NA NA 1 .
..
## $ has_company : Factor w/ 2 levels "False","True": NA NA NA NA NA NA NA NA NA NA NA ...
## $ has_gtin : Factor w/ 2 levels "False","True": NA NA NA NA NA NA NA NA NA NA NA ...
## $ average_stock : Factor w/ 6 levels "1-5","20-50",...: NA NA NA NA NA NA NA NA NA NA NA ...
## $ business_type : Factor w/ 3 levels "manufacturer",...: NA NA NA NA NA NA NA NA NA NA NA 3 ...
## $ declared_product_catalog_size: num NA NA NA NA NA NA NA NA NA NA NA ...
## $ declared_monthly_revenue : num NA NA NA NA NA NA NA NA NA NA 0 ...
## $ origin : Factor w/ 10 levels "direct_traffic",...: NA NA NA NA NA NA NA NA NA NA 7 .
..
## $ order_item_id : int 1 1 1 1 1 1 1 1 1 1 ...
## $ price : num 58.9 239.9 199 13 199.9 ...
## $ freight_value : num 13.3 19.9 17.9 12.8 18.1 ...
## $ order_status : Factor w/ 8 levels "approved","canceled",...: 4 4 4 4 4 4 4 4 4 4 ...
## $ review_score : int 5 4 5 4 5 4 4 5 1 4 ...
## $ review_comment_title : Factor w/ 4244 levels " BOM , MENOS LOGISTICA",...: NA NA NA NA NA NA NA NA
NA NA NA ...
## $ review_comment_message : Factor w/ 36508 levels "\n","\n\n","\n\n\n\n\n\n\n\n\n\n\n",...: 26004 NA
5280 NA 14858 NA NA NA 18476 NA ...
## $ EN_Review_comment_message : Factor w/ 35745 levels "\n\n\nOf any request made on the site, this was
the one that took to deliver !!!",...: 21991 NA 34468 NA 12758 NA NA NA 17455 NA ...
## $ conversion_time : num NA NA NA NA NA NA NA NA NA NA 0 ...
## $ delivery_time : num 7 16 8 6 25 7 8 5 10 2 ...
## $ feedback_time : num 8 17 9 7 26 8 9 6 11 3 ...
## $ delay_time : num -9 -3 -14 -6 -16 -15 -17 -16 0 -19 ...
## $ approval_time : num 0 0 0 0 0 2 0 1 1 0 ...
```

#Next, we're going to obtain the sentiment score from the variable "review comment in english" to keep this information as numeric

```
library(sentimentr)
library(stringr)
library(tidyverse)
library(tidytext)
library(tm)
```

```
## Loading required package: NLP
```

```
##
## Attaching package: 'NLP'
```

```
## The following object is masked from 'package:ggplot2':
##
## annotate
```

```
library(gmodels)
```

```
total4$EN_Review_comment_message <- as.character(total4$EN_Review_comment_message)
En_review = get_sentences(total4$EN_Review_comment_message)
df = sentiment_by(En_review)
total4$sentiment = df$save_sentiment
```

#Now we got the sentiment score. Let's look at the most popular words before dropping the text variables.

```
#install.packages("RColorBrewer")
library(wordcloud)
```

```
## Loading required package: RColorBrewer
```

```
library(RColorBrewer)
library(tidyverse)
library(tm)
library(SnowballC)

corpus = Corpus(VectorSource(total4$EN_Review_comment_message))
corpus[[1]][1]
```

```
## $content
## [1] "Perfect product delivered before combined."
```

```
#Conversion to Lowercase
corpus = tm_map(corpus, PlainTextDocument)
```

```
## Warning in tm_map.SimpleCorpus(corpus, PlainTextDocument): transformation
## drops documents
```

```
corpus = tm_map(corpus, tolower)
```

```
## Warning in tm_map.SimpleCorpus(corpus, tolower): transformation drops
## documents
```

```
#Removing Punctuation
corpus = tm_map(corpus, removePunctuation)
```

```
## Warning in tm_map.SimpleCorpus(corpus, removePunctuation): transformation
## drops documents
```

```
#Remove stopwords
corpus = tm_map(corpus, removeWords, c("cloth", stopwords("english")))
```

```
## Warning in tm_map.SimpleCorpus(corpus, removeWords, c("cloth",
## stopwords("english"))): transformation drops documents
```

```
# Stemming
corpus = tm_map(corpus, stemDocument)
```

```
## Warning in tm_map.SimpleCorpus(corpus, stemDocument): transformation drops
## documents
```

```
# Eliminate white spaces
corpus = tm_map(corpus, stripWhitespace)
```

```
## Warning in tm_map.SimpleCorpus(corpus, stripWhitespace): transformation
## drops documents
```

```
corpus[[1]][1]
```

```
## $content
## [1] "perfect product deliv combin"
```

```
#Next step is extracting the word frequencies, to be used as tags, for building the word cloud:
DTM <- TermDocumentMatrix(corpus)
mat <- as.matrix(DTM)
f <- sort(rowSums(mat), decreasing=TRUE)
dat <- data.frame(word = names(f), freq=f)
head(dat, 5)
```

```
##          word  freq
## product  product 22968
## receiv   receiv  8130
## good      good   7532
## deliveri deliveri 7253
## deliv     deliv  6491
```

```
set.seed(100)
wordcloud(words = dat$word, freq = dat$freq, min.freq = 3, max.words=250, random.order=FALSE, rot.per=0.30,
          colors=brewer.pal(8, "Dark2"))
```



#4. Looking at the structure of our data there are some variables classified as a factor with too many levels, that's because they should be classified as text. We are going to remove them since we don't need them for our analysis. We'll also remove other variables like the lead behaviour profile or the has_gtin, since we don't know the meaning of them.

```
remove04 = c("review_comment_message", "review_comment_title", "EN_Review_comment_message", "has_gtin", "lead_behaviour_profile")

total5 = total4 %>% dplyr::select(-remove04)
str(total5)
```

```
remove04 = c("review_comment_message", "review_comment_title", "EN_Review_comment_message", "has_gtin", "lead_behaviour_profile")

total5 = total4 %>% dplyr::select(-remove04)
str(total5)
```

```
total5 = total4 %>% dplyr::select(-remove04)
str(total5)
```

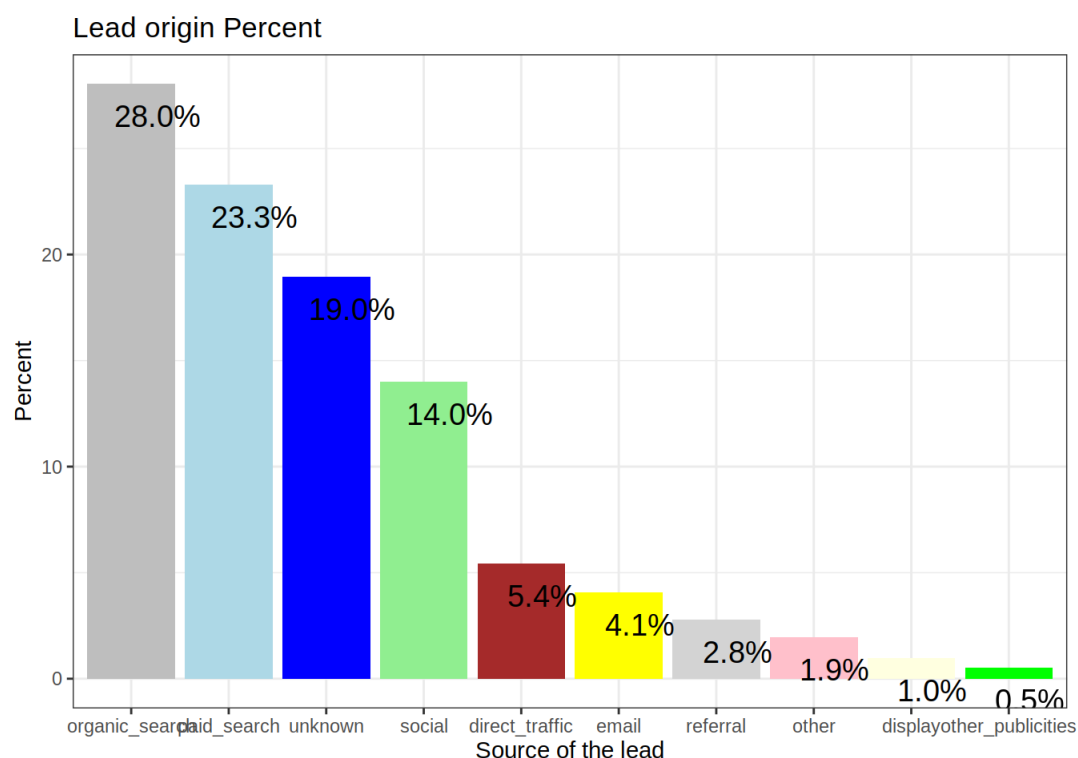
```
## 'data.frame': 121720 obs. of 19 variables:
## $ business_segment : Factor w/ 33 levels "air_conditioning",...: NA NA NA NA NA NA NA NA NA NA 6
...
## $ lead_type : Factor w/ 8 levels "industry","offline",...: NA NA NA NA NA NA NA NA NA NA 4 ...
## $ has_company : Factor w/ 2 levels "False","True": NA NA NA NA NA NA NA NA NA NA NA ...
## $ average_stock : Factor w/ 6 levels "1-5","20-50",...: NA NA NA NA NA NA NA NA NA NA NA ...
## $ business_type : Factor w/ 3 levels "manufacturer",...: NA NA NA NA NA NA NA NA NA NA NA 3 ...
## $ declared_product_catalog_size: num NA NA NA NA NA NA NA NA NA NA NA ...
## $ declared_monthly_revenue : num NA NA NA NA NA NA NA NA NA NA 0 ...
## $ origin : Factor w/ 10 levels "direct_traffic",...: NA NA NA NA NA NA NA NA NA NA 7 .
..
## $ order_item_id : int 1 1 1 1 1 1 1 1 1 1 ...
## $ price : num 58.9 239.9 199 13 199.9 ...
## $ freight_value : num 13.3 19.9 17.9 12.8 18.1 ...
## $ order_status : Factor w/ 8 levels "approved","canceled",...: 4 4 4 4 4 4 4 4 4 4 ...
## $ review_score : int 5 4 5 4 5 4 4 5 1 4 ...
## $ conversion_time : num NA NA NA NA NA NA NA NA NA NA 0 ...
## $ delivery_time : num 7 16 8 6 25 7 8 5 10 2 ...
## $ feedback_time : num 8 17 9 7 26 8 9 6 11 3 ...
## $ delay_time : num -9 -3 -14 -6 -16 -15 -17 -16 0 -19 ...
## $ approval_time : num 0 0 0 0 0 2 0 1 1 0 ...
## $ sentiment : num 0.335 0 0.769 0 0.158 ...
```

#5. Now let's understand better the distribution and relationship between our variables

#5.1 How many leads come from each origin source?

In this variable there are too many missing values, but before we drop it, let's see what we can find out from the observations we have omitting those missing values using `drop_na`

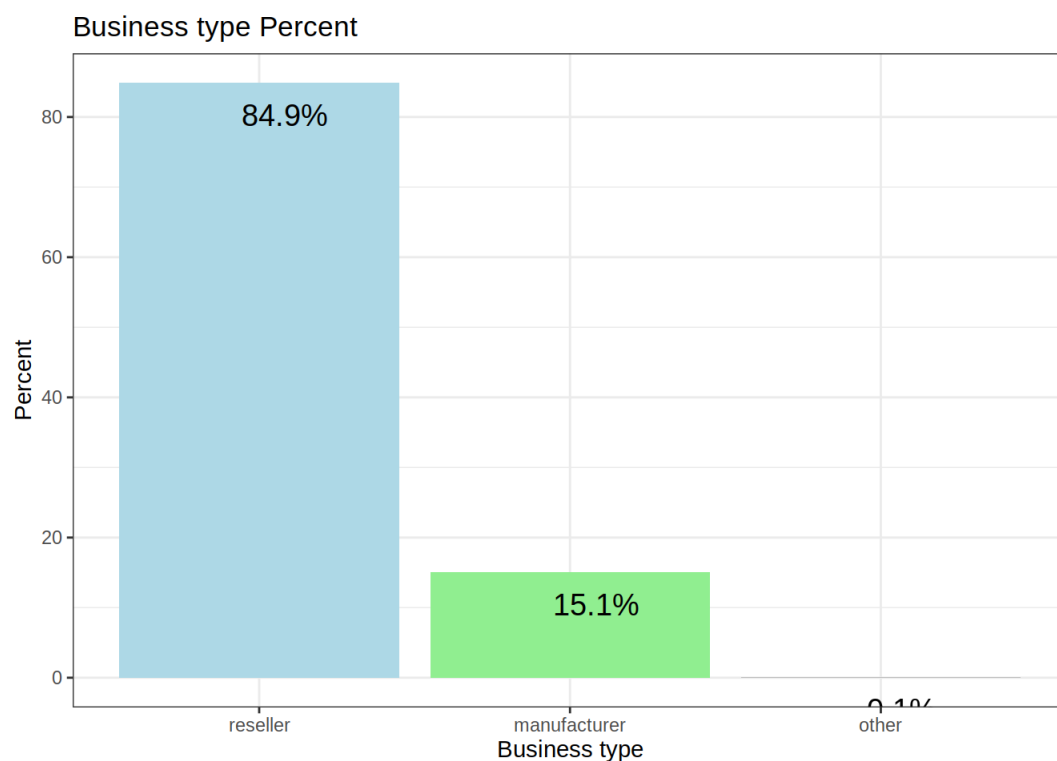
```
total5 %>%
  drop_na(origin) %>%
  group_by(origin) %>%
  summarise(Count = n()) %>%
  mutate(percent = prop.table(Count)*100) %>%
  ggplot(aes(reorder(origin, -percent), percent), fill = origin, na.rm = TRUE) +
  geom_col(fill = c("grey", "light blue", "blue", "light green", "brown", "yellow", "light grey", "pink", "light yellow", "green")) +
  geom_text(aes(label = sprintf("%.1f%%", percent)), hjust = 0.2, vjust = 2, size = 5) +
  theme_bw() +
  xlab("Source of the lead") + ylab("Percent") + ggtitle("Lead origin Percent")
```



```
#Results: Paid search and Organic search seemst to be the best way to gain more leads, followed by social
```

```
#5.2 Which business type is the most common?
```

```
total5 %>%
  drop_na(business_type) %>%
  group_by(business_type) %>%
  summarise(Count = n())%>%
  mutate(percent = prop.table(Count)*100)%>%
  ggplot(aes(reorder(business_type, -percent), percent), fill = business_type, na.rm = TRUE)+
  geom_col(fill = c("light blue", "light green", "grey"))+
  geom_text(aes(label = sprintf("%.1f%%", percent)), hjust = 0.2, vjust = 2, size = 5)+
  theme_bw()+
  xlab("Business type") + ylab("Percent") + ggtitle("Business type Percent")
```



```
#Results: 84% are resellers, so this is the type of business that should be targeted by Olist
```

```
#5.3 Can we group the different categories from the business segment into less categories?
```

```
#We can reduce the number of levels by grouping those with less frequency, but in this case it's better to group them using human judgement into similar categories
```

```
install.packages("ggplot2")
```

```
## Installing package into '/home/achaparro/R/x86_64-pc-linux-gnu-library/3.6'
## (as 'lib' is unspecified)
```

```
library(ggplot2)
library(ggpubr)
```


```
## Loading required package: magrittr
```

```
##
## Attaching package: 'magrittr'
```

```
## The following object is masked from 'package:purrr':
##
##   set_names
```

```
## The following object is masked from 'package:tidyr':
##
##   extract
```

```
ggplot(total5, aes(x = business_segment, na.rm = TRUE)) +
  geom_bar(fill = "#0073C2FF") +
  theme_pubclean()
```



```
library(dplyr)
df <- total5 %>%
  drop_na(business_segment) %>%
  group_by(business_segment) %>%
  summarise(counts = n())
df
```

```
## # A tibble: 33 x 2
##   business_segment      counts
##   <fct>              <int>
## 1 air_conditioning         6
## 2 audio_video_electronics 308
## 3 baby                    46
## 4 bags_backpacks          148
## 5 bed_bath_table          200
## 6 books                   111
## 7 car_accessories         211
## 8 computers                147
## 9 construction_tools_house_garden 356
## 10 fashion_accessories     73
## # ... with 23 more rows
```




```
total5$business_segment_new <- fct_collapse(total5$business_segment,
  home = c("home_appliances", "home_decor", "home_office_furniture", "household_utilities", "bed_bath_table",
  "construction_tools_house_garden", "air_conditioning"),
  food = c("food_supplement", "food_drink"),
  health_sport = c("sports_leisure", "health_beauty", "bags_backpacks"),
  electronics = c("audio_video_electronics", "computers", "games_consoles", "phone_mobile", "small_appliances", "watches"),
  free_time = c("gifts", "party", "fashion_accessories", "books", "handcrafted", "music_instruments"),
  kids = c("baby", "toys"),
  car = "car_accessories",
  pet = "pet",
  other = c("stationery", "religious"))

#ggplot(total5, aes(x = business_segment_new, na.rm = TRUE)) +
#  geom_bar(fill = "#0073C2FF", na.rm = TRUE) +
#  theme_pubclean()

?ggplot
total5 %>%
  drop_na(business_segment_new) %>%
  group_by(business_segment_new) %>%
  summarise(Count = n()) %>%
  mutate(percent = prop.table(Count)*100) %>%
  ggplot(aes(reorder(business_segment_new, -percent), percent), fill = business_type, na.rm = TRUE)+
  geom_col(fill = c("#66CC99", "#999999", "#CC6666", "#E69F00", "#56B4E9", "#009E73", "#F0E442", "#0072B2",
"#D55E00", "#CC79A7", "#9999CC"))+
  geom_text(aes(label = sprintf("%.1f%%", percent)), hjust = 0.2, vjust = 2, size = 5)+
  theme_bw()+
  xlab("Business segment new") + ylab("Percent") + ggtitle("Business segment Percent")
```

Business segment Percent



```
#Results: The segment more popular is Home, health&sport and electronics.

#Now we can drop the old "Business_segment" variable that contained too many levels

total5$business_segment <- NULL

str(total5)
```

```
## 'data.frame': 121720 obs. of 19 variables:
## $ lead_type : Factor w/ 8 levels "industry","offline",...: NA NA NA NA NA NA NA NA NA
4 ...
## $ has_company : Factor w/ 2 levels "False","True": NA NA NA NA NA NA NA NA NA ...
## $ average_stock : Factor w/ 6 levels "1-5","20-50",...: NA NA NA NA NA NA NA NA NA ...
## $ business_type : Factor w/ 3 levels "manufacturer",...: NA NA NA NA NA NA NA NA NA 3 ...
## $ declared_product_catalog_size: num NA NA NA NA NA NA NA NA NA NA ...
## $ declared_monthly_revenue : num NA NA NA NA NA NA NA NA NA 0 ...
## $ origin : Factor w/ 10 levels "direct_traffic",...: NA NA NA NA NA NA NA NA NA 7 .
..
## $ order_item_id : int 1 1 1 1 1 1 1 1 1 1 ...
## $ price : num 58.9 239.9 199 13 199.9 ...
## $ freight_value : num 13.3 19.9 17.9 12.8 18.1 ...
## $ order_status : Factor w/ 8 levels "approved","canceled",...: 4 4 4 4 4 4 4 4 4 ...
## $ review_score : int 5 4 5 4 5 4 4 5 1 4 ...
## $ conversion_time : num NA NA NA NA NA NA NA NA NA 0 ...
## $ delivery_time : num 7 16 8 6 25 7 8 5 10 2 ...
## $ feedback_time : num 8 17 9 7 26 8 9 6 11 3 ...
## $ delay_time : num -9 -3 -14 -6 -16 -15 -17 -16 0 -19 ...
## $ approval_time : num 0 0 0 0 0 2 0 1 1 0 ...
## $ sentiment : num 0.335 0 0.769 0 0.158 ...
## $ business_segment_new : Factor w/ 11 levels "home","electronics",...: NA NA NA NA NA NA NA NA NA
5 ...
```

```
summary(total5)
```

```
##          lead_type      has_company      average_stock
## online_medium: 2151   False:      29   1-5      :    10
## online_big    : 1924   True :      59   20-50   :     8
## online_small  : 501    NA's :121632   200+     :     8
## industry     : 361                5-20     :    22
## offline      : 219                50-200    :    39
## (Other)      : 305                unknown:     4
## NA's         :116259                NA's :121629
##          business_type  declared_product_catalog_size
## manufacturer: 830     Min.   :    1
## other        :    3     1st Qu.:   30
## reseller     : 4667    Median :  100
## NA's         :116220    Mean   :  233
##              3rd Qu.:  300
##              Max.    :2000
##              NA's    :121651
## declared_monthly_revenue      origin      order_item_id
## Min.   :    0      organic_search: 3534   Min.   : 1.000
## 1st Qu.:    0      paid_search   : 2934   1st Qu.: 1.000
## Median :    0      unknown       : 2390   Median : 1.000
## Mean   : 11209      social        : 1763   Mean   : 1.199
## 3rd Qu.:    0      direct_traffic: 686    3rd Qu.: 1.000
## Max.   :50000000      (Other)      : 1296   Max.   :21.000
## NA's   :116208      NA's         :109117   NA's   :8398
##          price      freight_value      order_status      review_score
## Min.   : 0.85   Min.   : 0.00   delivered :110848   Min.   :1
## 1st Qu.: 39.90   1st Qu.: 13.08   shipped    : 1197   1st Qu.:3
## Median : 74.90   Median : 16.26   canceled   : 711   Median :5
## Mean   : 120.48   Mean   : 19.98   unavailable: 612   Mean   :4
## 3rd Qu.: 134.90   3rd Qu.: 21.15   invoiced   : 366   3rd Qu.:5
## Max.   :6735.00   Max.   :409.68   (Other)    : 366   Max.   :5
## NA's   :8398     NA's   :8398     NA's       : 7620   NA's   :7620
## conversion_time  delivery_time  feedback_time  delay_time
## Min.   : -2.0   Min.   : 0.00   Min.   : -111.00   Min.   : -147.00
## 1st Qu.: 4.0   1st Qu.: 7.00   1st Qu.: 8.00     1st Qu.: -17.00
## Median : 11.0   Median : 10.00   Median : 11.00     Median : -13.00
## Mean   : 25.1   Mean   : 12.42   Mean   : 13.28     Mean   : -12.04
## 3rd Qu.: 23.0   3rd Qu.: 16.00   3rd Qu.: 17.00     3rd Qu.: -7.00
## Max.   :427.0   Max.   :210.00   Max.   : 148.00     Max.   : 188.00
## NA's   :116208   NA's   :10873   NA's       :7620   NA's   :10873
## approval_time    sentiment      business_segment_new
## Min.   : 0.00   Min.   : -1.62380   home       : 1901
## 1st Qu.: 0.00   1st Qu.: 0.00000   health_sport: 1207
## Median : 0.00   Median : 0.00000   electronics : 1197
## Mean   : 0.53   Mean   : 0.12190   pet         : 301
## 3rd Qu.: 1.00   3rd Qu.: 0.08866   free_time   : 251
## Max.   :188.00   Max.   : 2.64545   (Other)     : 654
## NA's   :7782                NA's       :116209
```

#5.4 We also may want to see the variances of our variables to find out if all of them will be adding useful information to our analysis. (first let's look at our numeric variables)

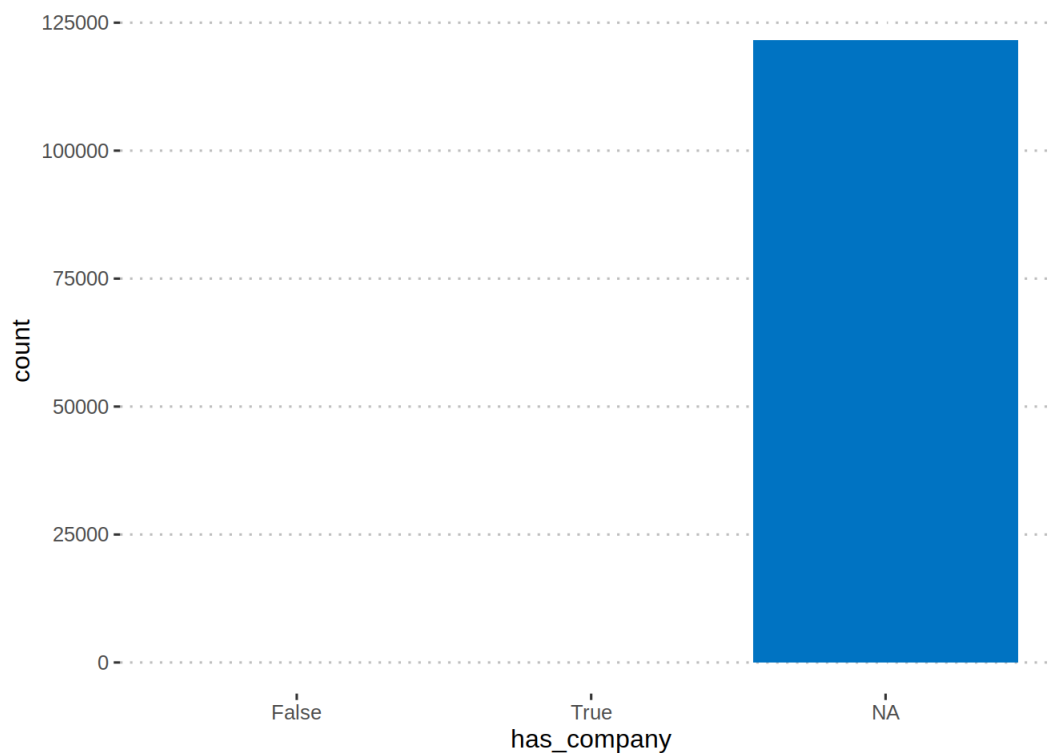
```
total_num <- total5[-c(1:4,7,11, 19)]
sapply(total_num, var, na.rm=TRUE)
```

```
## declared_product_catalog_size      declared_monthly_revenue
##          1.241721e+05                4.652727e+11
##          order_item_id              price
##          4.998388e-01                3.359069e+04
##          freight_value              review_score
##          2.491008e+02                1.994452e+00
##          conversion_time            delivery_time
##          2.176785e+03                8.928986e+01
##          feedback_time              delay_time
##          6.247500e+01                1.032310e+02
##          approval_time              sentiment
##          1.387445e+00                9.515489e-02
```


```
#The variable order_item_id revenue have a very low variance0, so we decide to drop it
total5$order_item_id <- NULL

#And for the categorical variables, we may want to have a look to the observations of each level

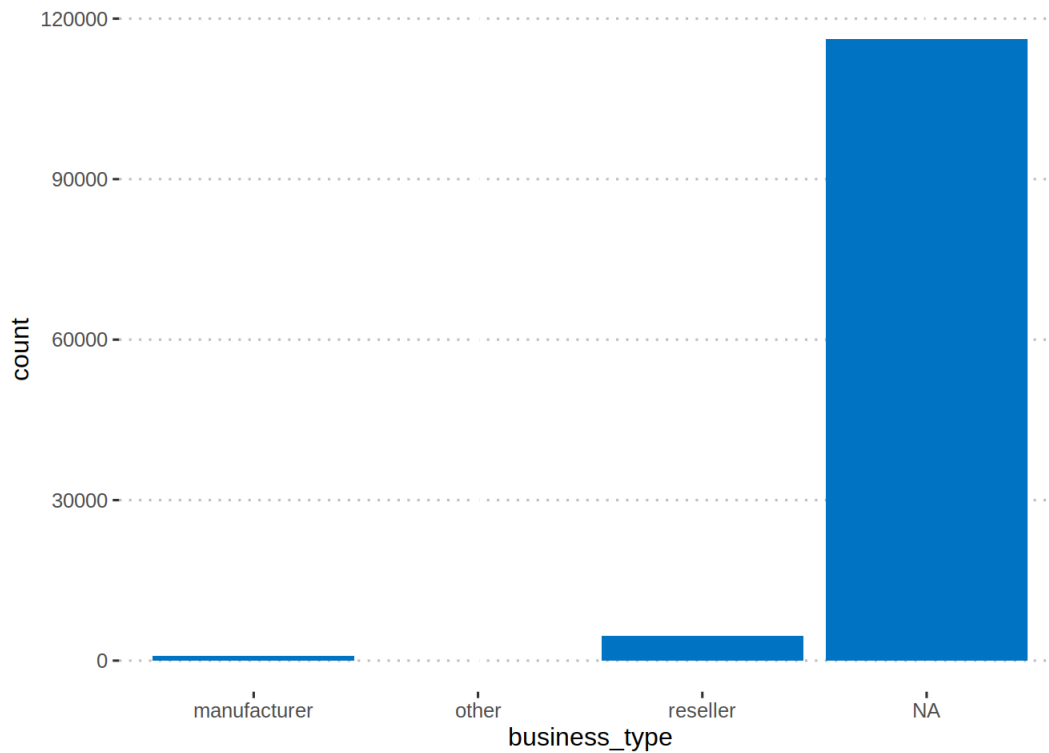
ggplot(total5, aes(has_company)) +
  geom_bar(fill = "#0073C2FF") +
  theme_pubclean()
```




```
ggplot(total5, aes(average_stock)) +
  geom_bar(fill = "#0073C2FF") +
  theme_pubclean()
```



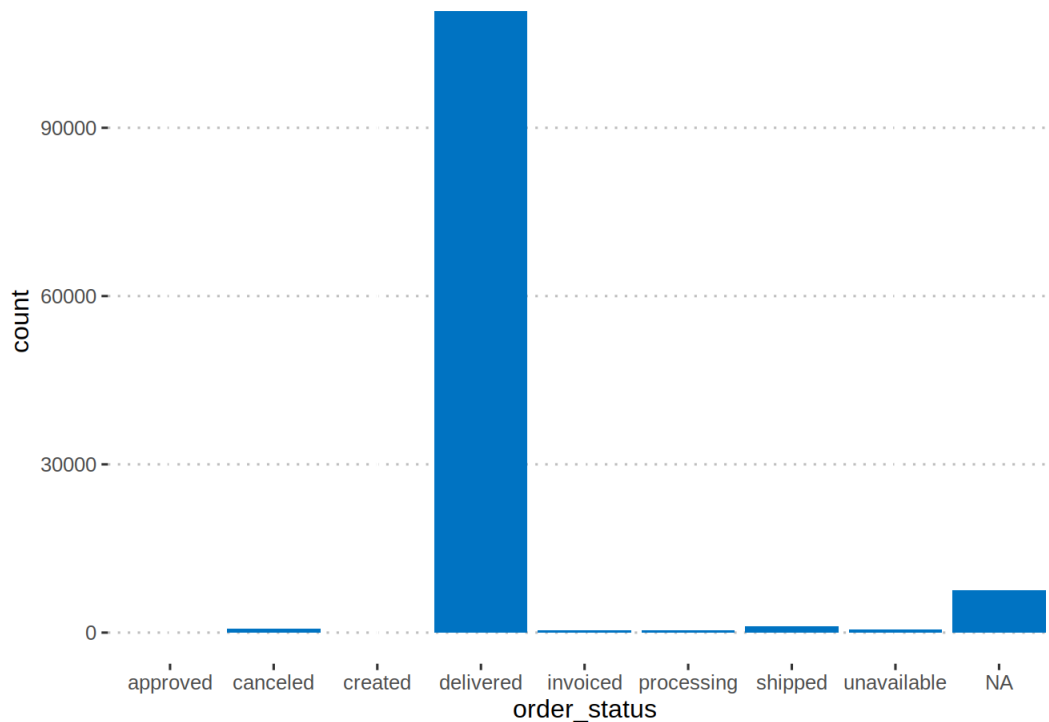
```
ggplot(total5, aes(business_type)) +
  geom_bar(fill = "#0073C2FF") +
  theme_pubclean()
```



```
ggplot(total5, aes(origin)) +
  geom_bar(fill = "#0073C2FF") +
  theme_pubclean()
```



```
ggplot(total5, aes(order_status)) +
  geom_bar(fill = "#0073C2FF") +
  theme_pubclean()
```



#Looking at the results, we decide to drop the variables: has_company, average_stock and order_status, since most of their observations are clasified in only one level, so they won't bring usefull information to our analysis

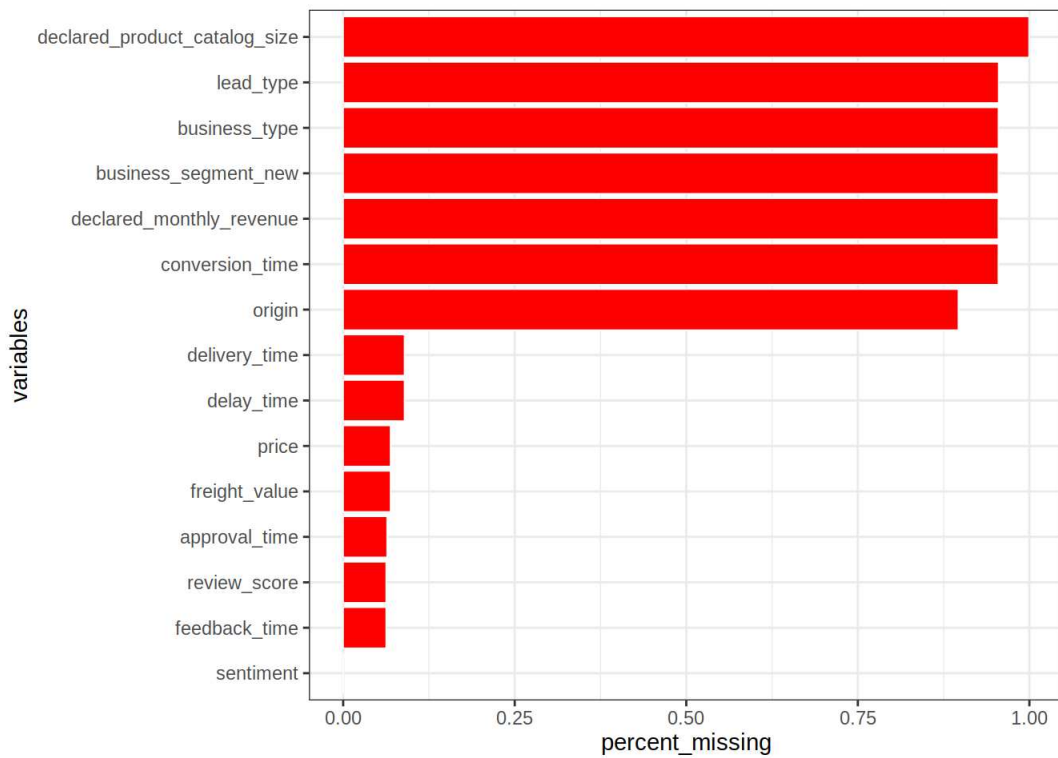
```
total5$has_company <- NULL
total5$average_stock <- NULL
total5$order_status <- NULL
```

#6.Handling the remaining missing values

```
missing_data <- total5 %>% summarise_all(funs(sum(is.na())/n()))
```

```
## Warning: funs() is soft deprecated as of dplyr 0.8.0
## Please use a list of either functions or lambdas:
##
##   # Simple named list:
##   list(mean = mean, median = median)
##
##   # Auto named with `tibble::lst()`:
##   tibble::lst(mean, median)
##
##   # Using lambdas
##   list(~ mean(., trim = .2), ~ median(., na.rm = TRUE))
## This warning is displayed once per session.
```

```
missing_data <- gather(missing_data, key = "variables", value = "percent_missing")
ggplot(missing_data, aes(x = reorder(variables, percent_missing), y = percent_missing)) +
  geom_bar(stat = "identity", fill = "red", aes(color = I('white')), size = 0.3)+
  xlab('variables')+
  coord_flip()+
  theme_bw()
```



```
summary(total5)
```

```
##          lead_type          business_type
## online_medium: 2151  manufacturer: 830
## online_big   : 1924  other       : 3
## online_small : 501   reseller    : 4667
## industry     : 361   NA's        :116220
## offline      : 219
## (Other)      : 305
## NA's         :116259
## declared_product_catalog_size declared_monthly_revenue
## Min.   : 1          Min.   : 0
## 1st Qu.: 30          1st Qu.: 0
## Median : 100         Median : 0
## Mean   : 233         Mean   : 11209
## 3rd Qu.: 300         3rd Qu.: 0
## Max.   :2000         Max.   :50000000
## NA's   :121651       NA's   :116208
##          origin          price          freight_value          review_score
## organic_search: 3534  Min.   : 0.85  Min.   : 0.00  Min.   :1
## paid_search   : 2934  1st Qu.: 39.90  1st Qu.: 13.08  1st Qu.:3
## unknown       : 2390  Median : 74.90  Median : 16.26  Median :5
## social        : 1763  Mean   : 120.48  Mean   : 19.98  Mean   :4
## direct_traffic: 686   3rd Qu.: 134.90  3rd Qu.: 21.15  3rd Qu.:5
## (Other)       : 1296  Max.   :6735.00  Max.   :409.68  Max.   :5
## NA's          :109117  NA's   :8398    NA's   :8398    NA's   :7620
## conversion_time delivery_time feedback_time delay_time
## Min.   : -2.0  Min.   : 0.00  Min.   : -111.00  Min.   : -147.00
## 1st Qu.: 4.0   1st Qu.: 7.00  1st Qu.: 8.00    1st Qu.: -17.00
## Median : 11.0  Median : 10.00  Median : 11.00    Median : -13.00
## Mean   : 25.1  Mean   : 12.42  Mean   : 13.28    Mean   : -12.04
## 3rd Qu.: 23.0  3rd Qu.: 16.00  3rd Qu.: 17.00    3rd Qu.: -7.00
## Max.   :427.0  Max.   :210.00  Max.   : 148.00    Max.   : 188.00
## NA's   :116208  NA's   :10873  NA's   :7620     NA's   :10873
## approval_time sentiment          business_segment_new
## Min.   : 0.00  Min.   : -1.62380  home       : 1901
## 1st Qu.: 0.00  1st Qu.: 0.00000  health_sport: 1207
## Median : 0.00  Median : 0.00000  electronics : 1197
## Mean   : 0.53  Mean   : 0.12190  pet         : 301
## 3rd Qu.: 1.00  3rd Qu.: 0.08866  free_time   : 251
## Max.   :188.00  Max.   : 2.64545  (Other)     : 654
## NA's   :7782    NA's   :116209
```

```
#We can see that still there are 7 variables with a high% of missing values (>100,000 obs). My decision is dropping those variables.
```

```
total5$lead_type <- NULL
total5$business_type <- NULL
total5$declared_product_catalog_size <- NULL
total5$declared_monthly_revenue <- NULL
total5$origin <- NULL
total5$conversion_time <- NULL
total5$business_segment_new <- NULL
```

```
summary(total5)
```

```
##      price      freight_value      review_score      delivery_time
## Min.   : 0.85    Min.   : 0.00    Min.   :1      Min.   : 0.00
## 1st Qu.: 39.90   1st Qu.: 13.08   1st Qu.:3      1st Qu.: 7.00
## Median : 74.90   Median : 16.26   Median :5      Median : 10.00
## Mean   : 120.48   Mean   : 19.98   Mean   :4      Mean   : 12.42
## 3rd Qu.: 134.90   3rd Qu.: 21.15   3rd Qu.:5      3rd Qu.: 16.00
## Max.   :6735.00   Max.   :409.68   Max.   :5      Max.   :210.00
## NA's   :8398     NA's   :8398     NA's   :7620   NA's   :10873
## feedback_time    delay_time    approval_time    sentiment
## Min.   : -111.00   Min.   : -147.00   Min.   : 0.00    Min.   : -1.62380
## 1st Qu.: 8.00     1st Qu.: -17.00   1st Qu.: 0.00    1st Qu.: 0.00000
## Median : 11.00    Median : -13.00   Median : 0.00    Median : 0.00000
## Mean   : 13.28    Mean   : -12.04   Mean   : 0.53    Mean   : 0.12190
## 3rd Qu.: 17.00    3rd Qu.: -7.00   3rd Qu.: 1.00    3rd Qu.: 0.08866
## Max.   : 148.00   Max.   : 188.00   Max.   :188.00   Max.   : 2.64545
## NA's   :7620     NA's   :10873     NA's   :7782
```

```
#A better solution is remove the rows that contains more than 50% of missing values
#dat[-which(rowMeans(is.na(dat)) > 0.5), ] is not working, why?
```

```
#Now we'll remove the missing values of the remaining variables
```

```
total6 <- na.omit(total5)
summary(total6)
```

```
##      price      freight_value      review_score      delivery_time
## Min.   : 0.85    Min.   : 0.00    Min.   :1.000    Min.   : 0.00
## 1st Qu.: 39.90   1st Qu.: 13.08   1st Qu.:4.000    1st Qu.: 7.00
## Median : 74.90   Median : 16.25   Median :5.000    Median : 10.00
## Mean   : 119.81   Mean   : 19.94   Mean   :4.066    Mean   : 12.42
## 3rd Qu.: 133.90   3rd Qu.: 21.15   3rd Qu.:5.000    3rd Qu.: 16.00
## Max.   :6735.00   Max.   :409.68   Max.   :5.000    Max.   :210.00
## feedback_time    delay_time    approval_time    sentiment
## Min.   : -77.00   Min.   : -147.00   Min.   : 0.0000   Min.   : -1.6238
## 1st Qu.: 8.00     1st Qu.: -17.00   1st Qu.: 0.0000   1st Qu.: 0.0000
## Median : 11.00    Median : -13.00   Median : 0.0000   Median : 0.0000
## Mean   : 12.89    Mean   : -12.03   Mean   : 0.5244   Mean   : 0.1343
## 3rd Qu.: 17.00    3rd Qu.: -7.00   3rd Qu.: 1.0000   3rd Qu.: 0.1595
## Max.   :112.00   Max.   : 188.00   Max.   :31.0000   Max.   : 2.6454
```


```
#7. Correlation of remaining variables (looking for multicollinearity)
```

```
cor(total6)
```



```
##           price freight_value review_score delivery_time
## price      1.000000000  0.4128919858  0.002442175   0.06262833
## freight_value 0.412891986  1.0000000000 -0.033009004   0.21480068
## review_score  0.002442175 -0.0330090041  1.000000000  -0.30462955
## delivery_time 0.062628330  0.2148006822 -0.304629550  1.000000000
## feedback_time 0.067919606  0.2524741414 -0.263903562  0.87454660
## delay_time   -0.003572271 -0.0399560099 -0.230318097  0.59704373
## approval_time 0.006680589  0.0262119652 -0.019522745  0.08547120
## sentiment    0.007755309 -0.0008628659  0.256464702  -0.09825000
##           feedback_time  delay_time approval_time  sentiment
## price      0.06791961 -0.003572271  0.006680589  0.0077553088
## freight_value 0.25247414 -0.039956010  0.026211965 -0.0008628659
## review_score -0.26390356 -0.230318097 -0.019522745  0.2564647018
## delivery_time 0.87454660  0.597043726  0.085471197 -0.0982499969
## feedback_time 1.000000000  0.387446372  0.113606449 -0.0912748999
## delay_time   0.38744637  1.000000000  0.046670743 -0.0877782852
## approval_time 0.11360645  0.046670743  1.000000000 -0.0017662492
## sentiment    -0.09127490 -0.087778285 -0.001766249  1.0000000000
```

```
heatmap(cor(total6), Rowv= NA, Colv = NA)
```



```
library(corrplot)
```

```
## corrplot 0.84 loaded
```

```
corrplot(cor(total6), method="number", type="lower")
```




#Results: Feedback time (the number of days that pass since the customer buy an item until he/she writes a review) and Delivery time are high correlated (0.87), but we'll keep both since their autocorrelation is not over 0.95.

#8. Inspecting Distribution of the remaining data


```
x0 <- total6$price
h0<-hist(x0, breaks=10, col="red", xlab="Price",
  main="Histogram with Normal Curve")
x0fit<-seq(min(x0),max(x0),length=40)
y0fit<-dnorm(x0fit,mean=mean(x0),sd=sd(x0))
y0fit <- y0fit*diff(h0$mids[1:2])*length(x0)
lines(x0fit, y0fit, col="blue", lwd=2)
```

Histogram with Normal Curve



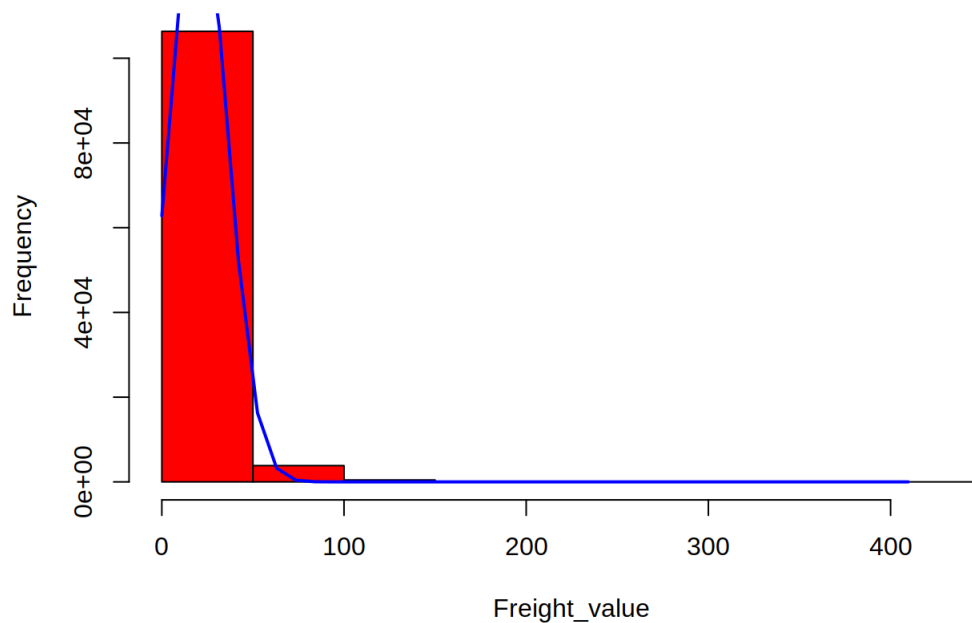
```
x1 <- total6$review_score
h1<-hist(x1, breaks=10, col="red", xlab="Review score",
  main="Histogram with Normal Curve")
x1fit<-seq(min(x1),max(x1),length=40)
yfit<-dnorm(x1fit,mean=mean(x1),sd=sd(x1))
yfit <- yfit*diff(h1$mids[1:2])*length(x1)
lines(x1fit, yfit, col="blue", lwd=2)
```

Histogram with Normal Curve



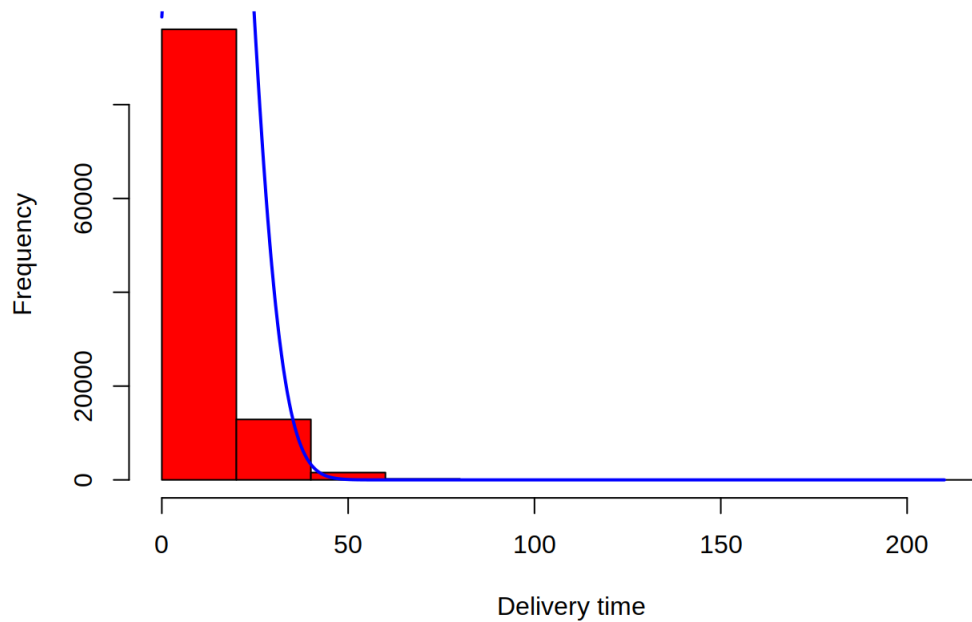
```
x <- total6$freight_value
h <-hist(x, breaks=10, col="red", xlab="Freight_value",
  main="Histogram with Normal Curve")
xfit<-seq(min(x),max(x),length=40)
yfit<-dnorm(xfit,mean=mean(x),sd=sd(x))
yfit <- yfit*diff(h$mids[1:2])*length(x)
lines(xfit, yfit, col="blue", lwd=2)
```

Histogram with Normal Curve



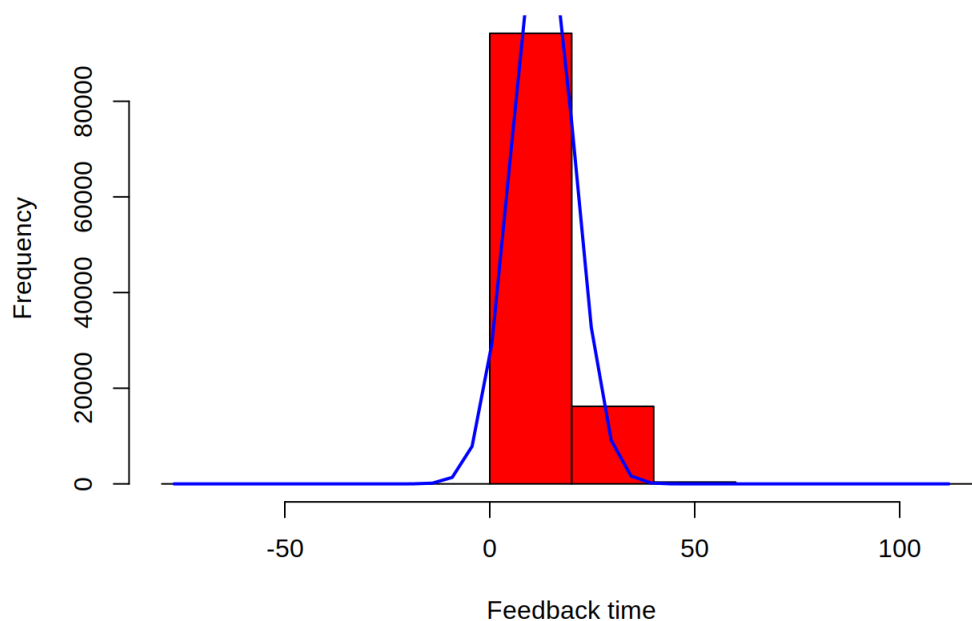
```
x1 <- total6$delivery_time
h1<-hist(x1, breaks=10, col="red", xlab="Delivery time",
  main="Histogram with Normal Curve")
x1fit<-seq(min(x1),max(x1),length=400)
y1fit<-dnorm(x1fit,mean=mean(x1),sd=sd(x1))
y1fit <- y1fit*diff(h$mids[1:2])*length(x1)
lines(x1fit, y1fit, col="blue", lwd=2)
```

Histogram with Normal Curve




```
x <- total6$feedback_time
h<-hist(x, breaks=10, col="red", xlab="Feedback time",
  main="Histogram with Normal Curve")
xfit<-seq(min(x),max(x),length=40)
yfit<-dnorm(xfit,mean=mean(x),sd=sd(x))
yfit <- yfit*diff(h$mids[1:2])*length(x)
lines(xfit, yfit, col="blue", lwd=2)
```

Histogram with Normal Curve




```
x <- total6$delay_time
h<-hist(x, breaks=10, col="red", xlab="Delay time",
  main="Histogram with Normal Curve")
xfit<-seq(min(x),max(x),length=40)
yfit<-dnorm(xfit,mean=mean(x),sd=sd(x))
yfit <- yfit*diff(h$mids[1:2])*length(x)
lines(xfit, yfit, col="blue", lwd=2)
```

Histogram with Normal Curve



```
x <- total6$approval
h<-hist(x, breaks=10, col="red", xlab="Approval time",
  main="Histogram with Normal Curve")
xfit<-seq(min(x),max(x),length=40)
yfit<-dnorm(xfit,mean=mean(x),sd=sd(x))
yfit <- yfit*diff(h$mids[1:2])*length(x)
lines(xfit, yfit, col="blue", lwd=2)
```

Histogram with Normal Curve



#3.Unsupervised analysis

```
#CLUSTER Analysis
```

```
library(cluster)
library(factoextra)
```

```
## Welcome! Related Books: `Practical Guide To Cluster Analysis in R` at https://goo.gl/13EFCZ
```

```
#install.packages("fastcluster")
library(fastcluster)
```

```
##
## Attaching package: 'fastcluster'
```

```
## The following object is masked from 'package:stats':
##
##      hclust
```

```
#km.out=kmeans(total6,2,nstart=5)
```

#When I try to run the cluster analysis I get this message: "Error: cannot allocate vector of size 45.8 Gb", so I'm going to use Spark to solve this memory issue

```
#library(sparklyr)
library(dplyr)
#sc = spark_connect(master = "local")
# We need to copy the data frame "total6" into the database "sc" as a table.
#total6_tbl = copy_to(sc, total6)
#src_tbls(sc)
```

```
#spark_kmeans <- ml_kmeans(total6_tbl, formula= NULL, k=3, max_iter = 10,
#features = c("price", "freight_value", "review_score", "delivery_time", "feedback_time", "delay_time", "approval_time", "sentiment"))
```

```
#summary(spark_kmeans)
```

```
#Time to compare the centers
# creating data frame from kmeans centers
#spark_kmeans_centers <- data.frame(spark_kmeans$centers)
# Printing centers of base and spark
#arrange(spark_kmeans_centers, review_score)
```


#Since I can't find a way to visualize the clusters using Spark, I'm going to drop randomly some observations so it will be possible to perform the cluster analysis from my laptop.

```
total6_reduced <- total6[sample(nrow(total6), 50000), ]
str(total6_reduced)
```

```
## 'data.frame': 50000 obs. of 8 variables:
## $ price : num 120 59.9 140 53.9 89.9 ...
## $ freight_value: num 20.3 17.7 15.7 19.4 16.4 ...
## $ review_score : int 1 5 5 5 4 5 5 5 5 5 ...
## $ delivery_time: num 8 15 17 13 12 18 10 7 4 5 ...
## $ feedback_time: num 9 16 18 14 13 19 11 8 5 6 ...
## $ delay_time : num -17 -5 -7 -31 -7 -16 -12 -17 -10 -10 ...
## $ approval_time: num 0 0 0 2 0 0 0 0 0 0 ...
## $ sentiment : num -0.158 0 0.694 0.411 0 ...
## - attr(*, "na.action")= 'omit' Named int 80 85 262 272 424 546 556 561 562 563 ...
## .. attr(*, "names")= chr "80" "85" "262" "272" ...
```


```
km.out=kmeans(total6_reduced,2,nstart=25)
#km.out$cluster
#?fviz_cluster
fviz_cluster(km.out, data= total6_reduced)
```

Cluster plot




```
plot(total6_reduced[,c("review_score", "sentiment")], col=(km.out$cluster+1),
     main="K-Means Clustering Results with K=2",
     xlab="Review score", ylab="Sentiment", pch=20, cex=2)
```

K-Means Clustering Results with K=2




```
plot(total16_reduced[,c("review_score","delivery_time")], col=(km.out$cluster+1),
     main="K-Means Clustering Results with K=2",
     xlab="Review score", ylab="delivery time", pch=20, cex=2)
```

K-Means Clustering Results with K=2



```
plot(total16_reduced[,c("review_score","delay_time")], col=(km.out$cluster+1),
     main="K-Means Clustering Results with K=2",
     xlab="Review score", ylab="Delay time", pch=20, cex=2)
```


K-Means Clustering Results with K=2



#SUPERVISED ANALYSIS (Regression predictive modeling) Treating the dependent variable as a numeric variable


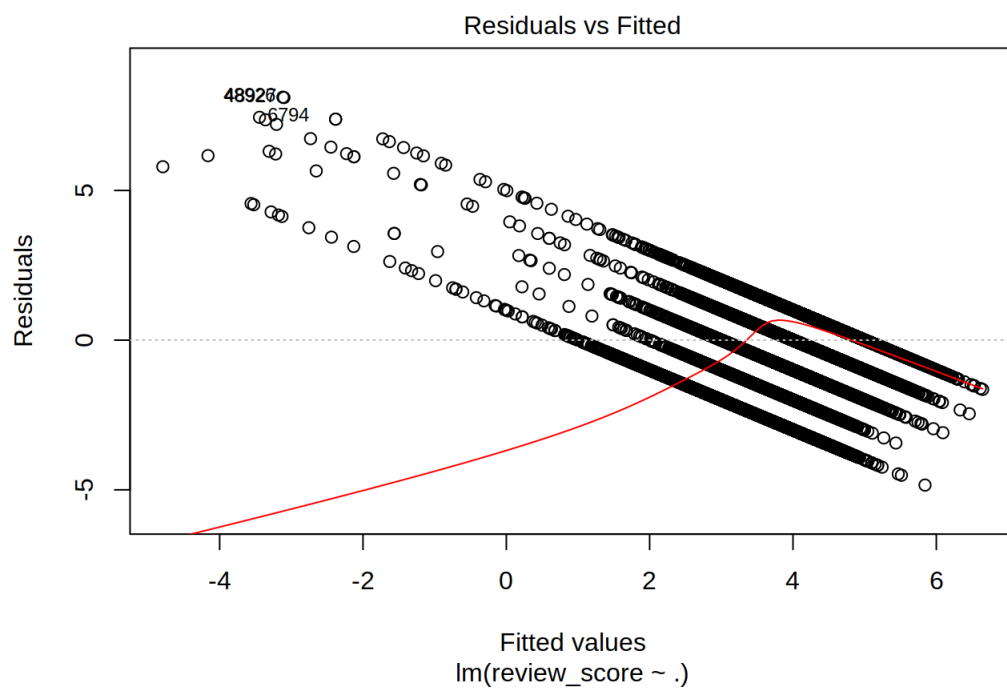
```
#Regression model
```



```
#Let's see which variables explain the variability of the review score and see if we find the best model to predict it
```

```
fit01=lm(review_score ~., data=total6)
summary(fit01)
```

```
##
## Call:
## lm(formula = review_score ~ ., data = total6)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.8422 -0.4930  0.4101  0.8853  8.1173
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  4.255e+00  1.167e-02  364.625 < 2e-16 ***
## price        7.474e-05  2.278e-05   3.280  0.00104 **
## freight_value 1.315e-03  2.770e-04   4.749  2.05e-06 ***
## delivery_time -3.208e-02  1.013e-03 -31.664 < 2e-16 ***
## feedback_time -5.127e-03  1.132e-03  -4.529  5.92e-06 ***
## delay_time    -8.861e-03  5.021e-04 -17.649 < 2e-16 ***
## approval_time  7.923e-03  3.872e-03   2.046  0.04074 *
## sentiment     9.654e-01  1.191e-02  81.073 < 2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.255 on 110824 degrees of freedom
## Multiple R-squared:  0.1479, Adjusted R-squared:  0.1479
## F-statistic: 2748 on 7 and 110824 DF, p-value: < 2.2e-16
```

```
plot(fit01)
```





```
#We found a model that explain the 14,79% of the review score variability.
#The most significant variables are the sentiment score(+), the delivery time(-), the delay time(-), the fee
dback time(-), the freight value (+). Followed by price, which is less significant and approval time, which
is not significant.

#We can run a 2nd regression model dropping that slightly significant variable and the R2 will not change.

fit02=lm(review_score ~ sentiment+delivery_time+delay_time+feedback_time+freight_value+price, data=total6)
summary(fit02)
```

```
##
## Call:
## lm(formula = review_score ~ sentiment + delivery_time + delay_time +
##      feedback_time + freight_value + price, data = total6)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -4.8390 -0.4937  0.4095  0.8863  8.1264
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)   4.257e+00  1.158e-02 367.539 < 2e-16 ***
## sentiment     9.656e-01  1.191e-02  81.092 < 2e-16 ***
## delivery_time -3.216e-02  1.012e-03 -31.762 < 2e-16 ***
## delay_time    -8.836e-03  5.019e-04 -17.605 < 2e-16 ***
## feedback_time -4.934e-03  1.128e-03  -4.373 1.22e-05 ***
## freight_value  1.316e-03  2.770e-04   4.752 2.02e-06 ***
## price         7.471e-05  2.278e-05   3.279 0.00104 **
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 1.255 on 110825 degrees of freedom
## Multiple R-squared:  0.1479, Adjusted R-squared:  0.1478
## F-statistic: 3206 on 6 and 110825 DF, p-value: < 2.2e-16
```

#Continue in part 2 because of memory issues

Carmen Marquez_Final project_part2

```
#Getting Total6 again:
```

```
library(dplyr)
```

```
##
```

```
## Attaching package: 'dplyr'
```

```
## The following objects are masked from 'package:stats':
```

```
##
```

```
## filter, lag
```

```
## The following objects are masked from 'package:base':
```

```
##
```

```
## intersect, setdiff, setequal, union
```

```
library(tidyverse)
```

```
## — Attaching packages —
```

```
———— tidyverse 1.2.1 —
```

```
## ✓ ggplot2 3.2.1      ✓ readr  1.3.1
```

```
## ✓ tibble  2.1.3      ✓ purrr  0.3.2
```

```
## ✓ tidyr   0.8.3      ✓ stringr 1.4.0
```

```
## ✓ ggplot2 3.2.1      ✓ forcats 0.4.0
```

```
## — Conflicts —
```

```
———— tidyverse_conflicts() —
```

```
## ✖ dplyr::filter() masks stats::filter()
```

```
## ✖ dplyr::lag() masks stats::lag()
```

```
library(sentimentr)
```

```
library(stringr)
```

```
library(tidyverse)
```

```
library(tidytext)
```

```
library(tm)
```

```
## Loading required package: NLP
```

```
##
```

```
## Attaching package: 'NLP'
```

```
## The following object is masked from 'package:ggplot2':
```

```
##
```

```
## annotate
```

```

library(gmodels)

setwd("/home/achaparro/personal/Carmen")
MQL <- read.csv("olist_marketing_qualified_leads_dataset.csv", na.strings = c("", "NA"))
closed_deals <- read.csv("olist_closed_deals_dataset.csv", na.strings = c("", "NA"))
total <- merge(closed_deals,MQL, by ="mql_id", all= TRUE)
order_items <- read.csv("olist_order_items_dataset.csv", na.strings = c("", "NA"))
order_reviews <- read.csv("olist_order_reviews_dataset_EN.csv", na.strings = c("", "NA"))
orders <- read.csv("olist_orders_dataset.csv", na.strings = c("", "NA"))
total1 <- merge(total,order_items, by ="seller_id", all= TRUE)
total2 <- merge(total1,orders, by ="order_id", all= TRUE)
total3 <- merge(total2,order_reviews, by ="order_id", all= TRUE)
remove02 = c("order_id", "seller_id" , "mql_id" , "sdr_id" , "sr_id","landing_page_id" , "product_id" , "customer_id" , "review_id")

total4 = total3 %>% dplyr::select(-remove02)
total4$won_date_new <- as.character(total4$won_date, format = "%Y-%m-%d")
total4$won_date_new <- as.Date(total4$won_date_new, format = "%Y-%m-%d")
total4$first_contact_date <- as.Date(total4$first_contact_date, format = "%Y-%m-%d")
total4$conversion_time <- (total4$won_date_new- total4$first_contact_date)
total4$conversion_time <- as.numeric(total4$conversion_time)

total4$order_delivered_customer_date <- as.Date(total4$order_delivered_customer_date, format = "%Y-%m-%d")
total4$order_purchase_timestamp <- as.Date(total4$order_purchase_timestamp, format = "%Y-%m-%d")
total4$delivery_time <- total4$order_delivered_customer_date - total4$order_purchase_timestamp
total4$delivery_time <- as.character(total4$delivery_time)
total4$delivery_time <- as.numeric(total4$delivery_time)
class(total4$delivery_time)

```

```
## [1] "numeric"
```

```

total4$review_creation_date <- as.Date(total4$review_creation_date, format = "%Y-%m-%d")
total4$feedback_time <- total4$review_creation_date - total4$order_purchase_timestamp
total4$feedback_time <- as.numeric(total4$feedback_time)

total4$order_estimated_delivery_date <- as.Date(total4$order_estimated_delivery_date, format = "%Y-%m-%d")
total4$delay_time <- total4$order_delivered_customer_date -
total4$order_estimated_delivery_date
total4$delay_time <- as.numeric(total4$delay_time)

total4$order_approved_at <- as.Date(total4$order_approved_at, format = "%Y-%m-%d")
total4$approval_time <- total4$order_approved_at - total4$order_purchase_timestamp
total4$approval_time <- as.numeric(total4$approval_time)

#3. We delete the date variables since we got the information we need from them in our new variables

remove03 = c("won_date", "shipping_limit_date", "order_purchase_timestamp", "order_approved_at", "order_delivered_carrier_date", "order_delivered_customer_date", "order_estimated_delivery_date", "review_creation_date", "review_answer_timestamp", "won_date_new", "first_contact_date")

total4 = total4 %>% dplyr::select(-remove03)
total4$EN_Review_comment_message <- as.character(total4$EN_Review_comment_message)
En_review = get_sentences(total4$EN_Review_comment_message)
df = sentiment_by(En_review)
total4$sentiment = df$sentiment
remove04 = c("review_comment_message", "review_comment_title", "EN_Review_comment_message", "has_gtin", "lead_behaviour_profile")

total5 = total4 %>% dplyr::select(-remove04)

total5$business_segment_new <- fct_collapse(total5$business_segment,
  home = c("home_appliances", "home_decor", "home_office_furniture", "household_utilities", "bed_bath_table", "construction_tools_house_garden", "air_conditioning"),
  food = c("food_supplement", "food_drink"),
  health_sport = c("sports_leisure", "health_beauty", "bags_backpacks"),
  electronics = c("audio_video_electronics", "computers", "games_consoles", "phone_mobile", "small_appliances", "watches"),
  free_time = c("gifts", "party", "fashion_accessories", "books", "handcrafted", "music_instruments"),
  kids = c("baby", "toys"),
  car = "car_accessories",
  pet = "pet",
  other = c("stationery", "religious"))
total5$business_segment <- NULL
total_num <- total5[-c(1:4, 7, 11, 19)]
sapply(total_num, var, na.rm=TRUE)

```

## declared_product_catalog_size	declared_monthly_revenue
## 1.241721e+05	4.652727e+11
## order_item_id	price
## 4.998388e-01	3.359069e+04
## freight_value	review_score
## 2.491008e+02	1.994452e+00
## conversion_time	delivery_time
## 2.176785e+03	8.928986e+01
## feedback_time	delay_time
## 6.247500e+01	1.032310e+02
## approval_time	sentiment
## 1.387445e+00	9.515489e-02

```
#The variable order_item_id revenue have a very low variance0, so we decide to drop it
total5$order_item_id <- NULL
total5$has_company <- NULL
total5$average_stock <- NULL
total5$order_status <- NULL

total5$lead_type <- NULL
total5$business_type <- NULL
total5$declared_product_catalog_size <- NULL
total5$declared_monthly_revenue <- NULL
total5$origin <- NULL
total5$conversion_time <- NULL
total5$business_segment_new <- NULL

total6 <- na.omit(total5)
```

#Supervised analysis (Ordered logistic regression model) Treating the dependent variable as an ordinal categorical variable

#We can't use the regular logistic model since our dependent variable is not binary, it is a categorical variable with several levels that are ordered, like a rank. So we are going to run an ordered logistic models and improve it using the Stepwise AIC method.

```
#install.packages("stargazer")
library(stargazer)
```

```
##
## Please cite as:
```

```
## Hlavac, Marek (2018). stargazer: Well-Formatted Regression and Summary Statistics Tables.
```

```
## R package version 5.2.2. https://CRAN.R-project.org/package=stargazer
```

```
library(MASS)
```

```
##
## Attaching package: 'MASS'
```

```
## The following object is masked from 'package:dplyr':
##
##      select
```

```
str(total6)
```

```
## 'data.frame':    110832 obs. of  8 variables:
## $ price          : num  58.9 239.9 199 13 199.9 ...
## $ freight_value: num  13.3 19.9 17.9 12.8 18.1 ...
## $ review_score  : int   5 4 5 4 5 4 4 5 1 4 ...
## $ delivery_time : num   7 16 8 6 25 7 8 5 10 2 ...
## $ feedback_time : num   8 17 9 7 26 8 9 6 11 3 ...
## $ delay_time    : num  -9 -3 -14 -6 -16 -15 -17 -16 0 -19 ...
## $ approval_time: num   0 0 0 0 0 2 0 1 1 0 ...
## $ sentiment     : num  0.335 0 0.769 0 0.158 ...
## - attr(*, "na.action")= 'omit' Named int   80 85 262 272 424 546 556 561 562 563 ...
## ..- attr(*, "names")= chr  "80" "85" "262" "272" ...
```

```
total6$review_score=as.factor(total6$review_score)

olm <- polr(review_score ~., data=total6, Hess=TRUE, method="logistic")
summary(olm)
```



```
## Call:
## polr(formula = review_score ~ ., data = total6, Hess = TRUE,
##       method = "logistic")
##
## Coefficients:
##               Value Std. Error t value
## price           0.000175  3.988e-05   4.389
## freight_value    0.001750  4.511e-04   3.880
## delivery_time   -0.069864  2.378e-03  -29.378
## feedback_time    0.016854  2.544e-03   6.624
## delay_time      -0.011458  8.187e-04  -13.995
## approval_time    0.007750  6.067e-03   1.277
## sentiment       1.806175  2.489e-02  72.571
##
## Intercepts:
##      Value      Std. Error t value
## 1|2    -2.4580      0.0206  -119.3753
## 2|3    -2.1199      0.0201  -105.6911
## 3|4    -1.4947      0.0194   -77.2070
## 4|5    -0.5045      0.0188  -26.8441
##
## Residual Deviance: 252049.53
## AIC: 252071.53
```

```
print(olm)
```

```
## Call:
## polr(formula = review_score ~ ., data = total6, Hess = TRUE,
##       method = "logistic")
##
## Coefficients:
##      price freight_value delivery_time feedback_time  delay_time
## 0.0001750439 0.0017503629 -0.0698643319 0.0168543491 -0.0114577769
## approval_time      sentiment
## 0.0077497715 1.8061750692
##
## Intercepts:
##      1|2      2|3      3|4      4|5
## -2.4580144 -2.1199184 -1.4946896 -0.5045025
##
## Residual Deviance: 252049.53
## AIC: 252071.53
```

```
#To get the pvalues we store the coefficient table, then calculate the p-values and combine back with the table:
(ctable <- coef(summary(olm)))
```

```
##               Value      Std. Error      t value
## price           0.0001750439 3.988336e-05   4.388894
## freight_value    0.0017503629 4.510701e-04   3.880468
## delivery_time   -0.0698643319 2.378093e-03  -29.378297
## feedback_time    0.0168543491 2.544338e-03   6.624256
## delay_time      -0.0114577769 8.187273e-04  -13.994619
## approval_time    0.0077497715 6.067146e-03   1.277334
## sentiment       1.8061750692 2.488846e-02  72.570771
## 1|2             -2.4580144125 2.059065e-02  -119.375299
## 2|3             -2.1199184064 2.005768e-02  -105.691087
## 3|4             -1.4946896116 1.935950e-02   -77.207033
## 4|5             -0.5045024555 1.879382e-02  -26.844063
```

```
p <- pnorm(abs(ctable[, "t value"]), lower.tail = FALSE) * 2
(ctable <- cbind(ctable, "p value" = p))
```

	Value	Std. Error	t value	p value
## price	0.0001750439	3.988336e-05	4.388894	1.139285e-05
## freight_value	0.0017503629	4.510701e-04	3.880468	1.042557e-04
## delivery_time	-0.0698643319	2.378093e-03	-29.378297	1.039998e-189
## feedback_time	0.0168543491	2.544338e-03	6.624256	3.490006e-11
## delay_time	-0.0114577769	8.187273e-04	-13.994619	1.681287e-44
## approval_time	0.0077497715	6.067146e-03	1.277334	2.014844e-01
## sentiment	1.8061750692	2.488846e-02	72.570771	0.000000e+00
## 1 2	-2.4580144125	2.059065e-02	-119.375299	0.000000e+00
## 2 3	-2.1199184064	2.005768e-02	-105.691087	0.000000e+00
## 3 4	-1.4946896116	1.935950e-02	-77.207033	0.000000e+00
## 4 5	-0.5045024555	1.879382e-02	-26.844063	9.894265e-159

#Interpretation of the 1st ordinal logistic model:

#Since the p-value for all the variables <0.05, hence they are statistically significant at 95% CI. The variable with the biggest pvalue is the approval time.

#As our predictive variables are continuous they can be interpreted as: E.g. With 1 unit increase in the delivery time the log of odds of a customer giving a better review score decreases by 0.069

#The intercepts can be interpreted in the following way: E.g. 1|2 means the log of odds of giving a review of 1, versus giving a review of 2,3,4 or 5.

#Using stepAIC to improve the model

```
step <- stepAIC(olm, direction="both")
```

```
## Start:  AIC=252071.5
## review_score ~ price + freight_value + delivery_time + feedback_time +
##   delay_time + approval_time + sentiment
##
##           Df    AIC
## - approval_time  1 252071
## <none>            252072
## - freight_value  1 252085
## - price          1 252091
## - feedback_time  1 252115
## - delay_time     1 252268
## - delivery_time  1 253080
## - sentiment      1 258662
##
## Step:  AIC=252071.2
## review_score ~ price + freight_value + delivery_time + feedback_time +
##   delay_time + sentiment
##
##           Df    AIC
## <none>            252071
## + approval_time  1 252072
## - freight_value  1 252084
## - price          1 252091
## - feedback_time  1 252116
## - delay_time     1 252267
## - delivery_time  1 253085
## - sentiment      1 258665
```

```
step$anova
```

```
## Stepwise Model Path
## Analysis of Deviance Table
##
## Initial Model:
## review_score ~ price + freight_value + delivery_time + feedback_time +
##   delay_time + approval_time + sentiment
##
## Final Model:
## review_score ~ price + freight_value + delivery_time + feedback_time +
##   delay_time + sentiment
##
##
##
##           Step Df Deviance Resid. Df Resid. Dev      AIC
## 1                110821    252049.5 252071.5
## 2 - approval_time  1 1.633586    110822    252051.2 252071.2
```

```
print(step)
```

```
## Call:
## polr(formula = review_score ~ price + freight_value + delivery_time +
##   feedback_time + delay_time + sentiment, data = total6, Hess = TRUE,
##   method = "logistic")
##
## Coefficients:
##           price freight_value delivery_time feedback_time   delay_time
## 0.0001748223  0.0017503217 -0.0699853660  0.0170874299 -0.0114288096
##      sentiment
## 1.8065426012
##
## Intercepts:
##           1|2           2|3           3|4           4|5
## -2.4608642 -2.1227753 -1.4975629 -0.5073952
##
## Residual Deviance: 252051.16
## AIC: 252071.16
```

#As we can see the variable approval time has been removed, and although the AIC has not improved too much, now the model is now more simple.

```
(ctable <- coef(summary(step)))
```

```
##           Value Std. Error t value
## price      0.0001748223 3.987916e-05  4.383802
## freight_value 0.0017503217 4.509722e-04  3.881219
## delivery_time -0.0699853660 2.376739e-03 -29.445967
## feedback_time 0.0170874299 2.538474e-03  6.731378
## delay_time   -0.0114288096 8.183619e-04 -13.965471
## sentiment    1.8065426012 2.488729e-02  72.588957
## 1|2          -2.4608641891 2.046992e-02 -120.218582
## 2|3          -2.1227752636 1.993296e-02 -106.495731
## 3|4          -1.4975629234 1.922876e-02 -77.881401
## 4|5          -0.5073951975 1.865700e-02 -27.195963
```

```
p1 <- pnorm(abs(ctable[, "t value"]), lower.tail = FALSE) * 2
(ctable <- cbind(ctable, "p value" = p1))
```

```
##           Value Std. Error t value p value
## price      0.0001748223 3.987916e-05  4.383802 1.166258e-05
## freight_value 0.0017503217 4.509722e-04  3.881219 1.039343e-04
## delivery_time -0.0699853660 2.376739e-03 -29.445967 1.417896e-190
## feedback_time 0.0170874299 2.538474e-03  6.731378 1.680638e-11
## delay_time   -0.0114288096 8.183619e-04 -13.965471 2.532245e-44
## sentiment    1.8065426012 2.488729e-02  72.588957 0.000000e+00
## 1|2          -2.4608641891 2.046992e-02 -120.218582 0.000000e+00
## 2|3          -2.1227752636 1.993296e-02 -106.495731 0.000000e+00
## 3|4          -1.4975629234 1.922876e-02 -77.881401 0.000000e+00
## 4|5          -0.5073951975 1.865700e-02 -27.195963 7.249713e-163
```

```

#Interpretation of the improved ordinal logistic model:
#All the variables are statistically significant at 95% CI. The variables with the biggest pvalue are now the freight value and the price.

#Our predictive variables are continuous so they can be interpreted as: E.g. With 1 unit increase in the delivery time the log of odds of a customer giving a better review score decreases by 0.069

#The intercepts can be interpreted in the following way: E.g. 1|2 means the log of odds of giving a review of 1, versus giving a review of 2,3,4 or 5.

#Test and train our logistic model
#Set Testing Criteria -70/30
numberofobs = round(length(total6$review_score) * .7)

#Split Test and Train data
train <- total6[1:numberofobs,]
test <- total6[-(1:numberofobs),]

#Make predictions(Step)
setup2 <- test
setup2[, c("pred.prob")] <- predict(step, newdata=setup2, type="probs")
setup2[, c("pred.prob")] <- predict(step, newdata=setup2, type="class")
setup2$residuals <- residuals(step, type="response")

#Step AIC model confusion matrix
library(caret)

```

```
## Loading required package: lattice
```

```
##
## Attaching package: 'caret'
```

```
## The following object is masked from 'package:purrr':
##
## lift
```

```
confusionMatrix(setup2$pred.prob, test$review_score, positive="TRUE")
```

```
## Confusion Matrix and Statistics
##
##           Reference
## Prediction    1    2    3    4    5
##      1 1010  170  186  142  191
##      2    0    0    0    0    0
##      3    0    0    0    0    0
##      4    0    0    0    0    0
##      5 2760  983 2603 6176 19029
##
## Overall Statistics
##
##           Accuracy : 0.6027
##           95% CI : (0.5974, 0.6079)
##      No Information Rate : 0.578
##      P-Value [Acc > NIR] : < 2.2e-16
##
##           Kappa : 0.1085
##
##  McNemar's Test P-Value : NA
##
## Statistics by Class:
##
##           Class: 1 Class: 2 Class: 3 Class: 4 Class: 5
## Sensitivity      0.26790 0.00000 0.00000    0.00 0.9901
## Specificity      0.97663 1.00000 1.00000    1.00 0.1075
## Pos Pred Value   0.59447      NaN      NaN      NaN 0.6031
## Neg Pred Value   0.91252 0.96532 0.91612    0.81 0.8876
## Prevalence       0.11338 0.03468 0.08388    0.19 0.5780
## Detection Rate   0.03038 0.00000 0.00000    0.00 0.5723
## Detection Prevalence 0.05110 0.00000 0.00000    0.00 0.9489
## Balanced Accuracy 0.62227 0.50000 0.50000    0.50 0.5488
```

#Supervised analysis (Machine learning predictive modeling)

```
#Let's see which is the best model to predict the review score.
```

```
library(rattle)
```

```
## Rattle: A free graphical interface for data science with R.
## Version 5.2.0 Copyright (c) 2006-2018 Togaware Pty Ltd.
## Type 'rattle()' to shake, rattle, and roll your data.
```

```
library(DMwR)
```

```
## Loading required package: grid
```

```
## Registered S3 method overwritten by 'xts':
##   method      from
## as.zoo.xts zoo
```

```
## Registered S3 method overwritten by 'quantmod':
##   method      from
## as.zoo.data.frame zoo
```

```
library(caret)
library(lattice)
library(e1071)
library(tidyverse)

# 10-fold Cross-Validation
control <- trainControl(method="cv", number=10)
metric <- "Accuracy"

# Linear Discriminant Analysis (LDA)
set.seed(99)
fit.lda <- train(review_score ~., data=total6, method="lda", metric=metric, trControl=control)

# Classification and Regression Trees (CART)
set.seed(99)
fit.cart <- train(review_score~., data=total6, method="rpart", metric=metric, trControl=control)

# k-Nearest Neighbors (KNN)
set.seed(99)
fit.knn <- train(review_score~., data=total6, method="knn", metric=metric, trControl=control)

# Bayesian Generalized Linear Model - Logistic Regression
set.seed(99)
fit.logi <- train(review_score~., data=total6, method="bayesglm", metric=metric, trControl=control)

# Random Forest
set.seed(99)
fit.rf <- train(review_score~., data=total6, method="rf", metric=metric, trControl=control)

# Gradient Boosting Machines/XGBoost-Linear Model
set.seed(99)
fit.xgb <- train(review_score~., data=total6, method="xgbLinear", metric=metric, trControl=control)

# Gradient Boosting Machines/XGBoost-Tree Model
#set.seed(99)
#fit.xgb.t <- train(review_score~., data=total6, method="xgbTree", metric=metric, trControl=control)

# Select Best Model
# summarize accuracy of models
results <- resamples(list(lda=fit.lda, cart=fit.cart, knn=fit.knn, logi=fit.logi, rf=fit.rf, xgb.l=fit.xgb))
summary(results)
```

```
##
## Call:
## summary.resamples(object = results)
##
## Models: lda, cart, knn, logi, rf, xgb.l
## Number of resamples: 10
##
## Accuracy
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## lda      0.59379229 0.59524508 0.59683310 0.59685836 0.59763140 0.60010827
## cart     0.62501128 0.62677912 0.62866552 0.63012486 0.63280177 0.63758571
## knn      0.55724984 0.56071096 0.56317949 0.56274362 0.56502909 0.56677495
## logi     0.04736557 0.04806478 0.04876619 0.04887583 0.04933566 0.05143476
## rf       0.67421508 0.68026891 0.68319422 0.68168943 0.68323710 0.68528377
## xgb.l    0.64395922 0.64459081 0.64779392 0.64699730 0.64841089 0.65114600
##           NA's
## lda      0
## cart     0
## knn      0
## logi     0
## rf       0
## xgb.l    0
##
## Kappa
##           Min.      1st Qu.      Median      Mean      3rd Qu.      Max.
## lda      0.10346243 0.10736651 0.10920773 0.11100794 0.11303876 0.12144479
## cart     0.21658049 0.22326337 0.22858232 0.24102989 0.26525950 0.27279618
## knn      0.13439868 0.14378372 0.14639809 0.14592000 0.15109249 0.15257394
## logi     0.01147314 0.01223018 0.01293248 0.01302151 0.01350773 0.01552172
## rf       0.39918976 0.40638249 0.41238886 0.41040655 0.41559723 0.41663931
## xgb.l    0.28490837 0.28977149 0.29356106 0.29305427 0.29650678 0.30221779
##           NA's
## lda      0
## cart     0
## knn      0
## logi     0
## rf       0
## xgb.l    0
```

```
#The best model is Random Forest with a kappa of 0.40
```

```
# Summarize the Best Model
print(fit.rf)
```

```
## Random Forest
##
## 110832 samples
##      7 predictor
##      5 classes: '1', '2', '3', '4', '5'
##
## No pre-processing
## Resampling: Cross-Validated (10 fold)
## Summary of sample sizes: 99749, 99748, 99748, 99749, 99750, 99748, ...
## Resampling results across tuning parameters:
##
##      mtry Accuracy      Kappa
##      2      0.6802187 0.3766649
##      4      0.6816894 0.4104066
##      7      0.6775841 0.4070959
##
## Accuracy was used to select the optimal model using the largest value.
## The final value used for the model was mtry = 4.
```

```
summary(fit.rf)
```

##	Length	Class	Mode
## call	4	-none-	call
## type	1	-none-	character
## predicted	110832	factor	numeric
## err.rate	3000	-none-	numeric
## confusion	30	-none-	numeric
## votes	554160	matrix	numeric
## oob.times	110832	-none-	numeric
## classes	5	-none-	character
## importance	7	-none-	numeric
## importanceSD	0	-none-	NULL
## localImportance	0	-none-	NULL
## proximity	0	-none-	NULL
## ntree	1	-none-	numeric
## mtry	1	-none-	numeric
## forest	14	-none-	list
## y	110832	factor	numeric
## test	0	-none-	NULL
## inbag	0	-none-	NULL
## xNames	7	-none-	character
## problemType	1	-none-	character
## tuneValue	1	data.frame	list
## obsLevels	5	-none-	character
## param	0	-none-	list

Random Forest model appears to be the best choice machine learning model when we treat our dependant variable as a categorical one. We can see it's kappa it's the highest in comparison with the rest of the machine learning models.

However, I will choose as the best model for my goal the Ordinal Logistic Regression. The reason for making this choice is that although its accuracy level is not has high as the one we find in the rf model(0.60 Accuracy and Kappa 0.10), it is a good alternative model for interpreting which factors influence in the review score. Random forest model is more complex and less easy for interpretation.